

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_excel('oil_spill.xlsx')
df.head()
```

```
Out[2]:
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_10	...	f_41	f_42	f_43	f_44	f_45	f_46	f_47	f_48	f_49	target
0	1	2558	1506.09	456.63	90	6395000	40.88	7.89	29780.0	0.19	...	2850.00	1000.00	763.16	135.46	3.73	0	33243.19	65.74	7.95	1
1	2	22325	79.11	841.03	180	55812500	51.11	1.21	61900.0	0.02	...	5750.00	11500.00	9593.48	1648.80	0.60	0	51572.04	65.73	6.26	0
2	3	115	1449.85	608.43	88	287500	40.42	7.34	3340.0	0.18	...	1400.00	250.00	150.00	45.13	9.33	1	31692.84	65.81	7.84	1
3	4	1201	1562.53	295.65	66	3002500	42.40	7.97	18030.0	0.19	...	6041.52	761.58	453.21	144.97	13.33	1	37696.21	65.67	8.07	1
4	5	312	950.27	440.86	37	780000	41.43	7.03	3350.0	0.17	...	1320.04	710.63	512.54	109.16	2.58	0	29038.17	65.66	7.35	0

5 rows × 50 columns

```
In [3]: df.shape
```

```
Out[3]: (937, 50)
```

```
In [4]: df.columns
```

```
Out[4]: Index(['f_1', 'f_2', 'f_3', 'f_4', 'f_5', 'f_6', 'f_7', 'f_8', 'f_9', 'f_10',
              'f_11', 'f_12', 'f_13', 'f_14', 'f_15', 'f_16', 'f_17', 'f_18', 'f_19',
              'f_20', 'f_21', 'f_22', 'f_23', 'f_24', 'f_25', 'f_26', 'f_27', 'f_28',
              'f_29', 'f_30', 'f_31', 'f_32', 'f_33', 'f_34', 'f_35', 'f_36', 'f_37',
              'f_38', 'f_39', 'f_40', 'f_41', 'f_42', 'f_43', 'f_44', 'f_45', 'f_46',
              'f_47', 'f_48', 'f_49', 'target'],
              dtype='object')
```

```
In [5]: df.isnull().sum()
```

```
Out[5]: f_1      0
        f_2      0
        f_3      0
        f_4      0
        f_5      0
        f_6      0
        f_7      0
        f_8      0
        f_9      0
        f_10     0
        f_11     0
        f_12     0
        f_13     0
        f_14     0
        f_15     0
        f_16     0
        f_17     0
        f_18     0
        f_19     0
        f_20     0
        f_21     0
        f_22     0
        f_23     0
        f_24     0
        f_25     0
        f_26     0
        f_27     0
        f_28     0
        f_29     0
        f_30     0
        f_31     0
        f_32     0
        f_33     0
        f_34     0
        f_35     0
        f_36     0
        f_37     0
        f_38     0
        f_39     0
        f_40     0
        f_41     0
        f_42     0
        f_43     0
        f_44     0
        f_45     0
        f_46     0
        f_47     0
        f_48     0
        f_49     0
        target    0
        dtype: int64
```

```
In [6]: df.duplicated().sum()
```

```
Out[6]: 0
```

```
In [7]: df.dtypes
```

```
Out[7]: f_1      int64
        f_2      int64
        f_3     float64
        f_4     float64
        f_5      int64
        f_6      int64
        f_7     float64
        f_8     float64
        f_9     float64
        f_10    float64
        f_11    float64
        f_12    float64
        f_13    float64
        f_14    float64
        f_15    float64
        f_16    float64
        f_17    float64
        f_18    float64
        f_19    float64
        f_20    float64
        f_21    float64
        f_22    float64
        f_23     int64
        f_24    float64
        f_25    float64
        f_26    float64
        f_27    float64
        f_28    float64
        f_29    float64
        f_30    float64
        f_31    float64
        f_32    float64
        f_33    float64
        f_34    float64
        f_35     int64
        f_36     int64
        f_37    float64
        f_38    float64
        f_39     int64
        f_40     int64
        f_41    float64
        f_42    float64
        f_43    float64
        f_44    float64
        f_45    float64
        f_46     int64
        f_47    float64
        f_48    float64
        f_49    float64
        target   int64
dtype: object
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 937 entries, 0 to 936
Data columns (total 50 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   f_1         937 non-null    int64
1   f_2         937 non-null    int64
2   f_3         937 non-null    float64
3   f_4         937 non-null    float64
4   f_5         937 non-null    int64
5   f_6         937 non-null    int64
6   f_7         937 non-null    float64
7   f_8         937 non-null    float64
8   f_9         937 non-null    float64
9   f_10        937 non-null    float64
10  f_11        937 non-null    float64
11  f_12        937 non-null    float64
12  f_13        937 non-null    float64
13  f_14        937 non-null    float64
14  f_15        937 non-null    float64
15  f_16        937 non-null    float64
16  f_17        937 non-null    float64
17  f_18        937 non-null    float64
18  f_19        937 non-null    float64
19  f_20        937 non-null    float64
20  f_21        937 non-null    float64
21  f_22        937 non-null    float64
22  f_23        937 non-null    int64
23  f_24        937 non-null    float64
24  f_25        937 non-null    float64
25  f_26        937 non-null    float64
26  f_27        937 non-null    float64
27  f_28        937 non-null    float64
28  f_29        937 non-null    float64
29  f_30        937 non-null    float64
30  f_31        937 non-null    float64
31  f_32        937 non-null    float64
32  f_33        937 non-null    float64
33  f_34        937 non-null    float64
34  f_35        937 non-null    int64
35  f_36        937 non-null    int64
36  f_37        937 non-null    float64
37  f_38        937 non-null    float64
38  f_39        937 non-null    int64
39  f_40        937 non-null    int64
40  f_41        937 non-null    float64
41  f_42        937 non-null    float64
42  f_43        937 non-null    float64
43  f_44        937 non-null    float64
44  f_45        937 non-null    float64
45  f_46        937 non-null    int64
46  f_47        937 non-null    float64
47  f_48        937 non-null    float64
48  f_49        937 non-null    float64
49  target      937 non-null    int64
dtypes: float64(39), int64(11)
memory usage: 366.1 KB
```

```
In [9]: df.describe()
```

Out[9]:


	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_10	...	f_41
count	937.000000	937.000000	937.000000	937.000000	937.000000	9.370000e+02	937.000000	937.000000	937.000000	937.000000	...	937.000000
mean	81.588047	332.842049	698.707086	870.992209	84.121665	7.696964e+05	43.242721	9.127887	3940.712914	0.221003	...	933.928677
std	64.976730	1931.938570	599.965577	522.799325	45.361771	3.831151e+06	12.718404	3.588878	8167.427625	0.090316	...	1001.681331
min	1.000000	10.000000	1.920000	1.000000	0.000000	7.031200e+04	21.240000	0.830000	667.000000	0.020000	...	0.000000
25%	31.000000	20.000000	85.270000	444.200000	54.000000	1.250000e+05	33.650000	6.750000	1371.000000	0.160000	...	450.000000
50%	64.000000	65.000000	704.370000	761.280000	73.000000	1.863000e+05	39.970000	8.200000	2090.000000	0.200000	...	685.420000
75%	124.000000	132.000000	1223.480000	1260.370000	117.000000	3.304680e+05	52.420000	10.760000	3435.000000	0.260000	...	1053.420000
max	352.000000	32389.000000	1893.080000	2724.570000	180.000000	7.131500e+07	82.640000	24.690000	160740.000000	0.740000	...	11949.330000

8 rows × 50 columns

```
In [10]: df['target'].value_counts()
```

```
Out[10]: 0    896  
         1     41  
         Name: target, dtype: int64
```


```
In [11]: plt.figure(figsize= (9,7))  
sns.countplot(x = df['target'])  
plt.show()
```




```
In [12]: df.columns
```

```
Out[12]: Index(['f_1', 'f_2', 'f_3', 'f_4', 'f_5', 'f_6', 'f_7', 'f_8', 'f_9', 'f_10',  
              'f_11', 'f_12', 'f_13', 'f_14', 'f_15', 'f_16', 'f_17', 'f_18', 'f_19',  
              'f_20', 'f_21', 'f_22', 'f_23', 'f_24', 'f_25', 'f_26', 'f_27', 'f_28',  
              'f_29', 'f_30', 'f_31', 'f_32', 'f_33', 'f_34', 'f_35', 'f_36', 'f_37',  
              'f_38', 'f_39', 'f_40', 'f_41', 'f_42', 'f_43', 'f_44', 'f_45', 'f_46',  
              'f_47', 'f_48', 'f_49', 'target'],  
              dtype='object')
```

```
In [13]: for i in df.columns:  
sns.displot(df[i])  
plt.title(f'displot for {i}')  
plt.show()
```



```
In [14]: for i in df.columns:
sns.boxplot(df[i])
plt.title(f'Boxplot for {i}')
plt.show()
```



```
In [ ]:
```

Outlier to be treated

f\_8,

f\_10

f\_13


f\_41

f\_44

f\_47

```
In [15]: sns.boxplot(df['f_8'])
plt.title('boxplot for f_8')
```

```
Out[15]: Text(0.5, 1.0, 'boxplot for f_8')
```



```
In [16]: q1 = np.percentile(df['f_8'],25)
q3 = np.percentile(df['f_8'],75)
iqr = q3-q1
v_max = q3 + 1.5*iqr
print(v_max)
```


```
16.775
```

```
In [17]: print(df[df['f_8']>16.775].shape)
```

```
(38, 50)
```


```
In [18]: df['f_8'] = np.where(df['f_8']>16.775,16.775,df['f_8'])
sns.boxplot(x = df['f_8'])
```

```
Out[18]: <AxesSubplot:xlabel='f_8'>
```



```
In [19]: sns.boxplot(df['f_10'])
plt.title('boxplot for f_10')
```

```
Out[19]: Text(0.5, 1.0, 'boxplot for f_10')
```



```
In [20]: q1 = np.percentile(df['f_10'],25)
q3 = np.percentile(df['f_10'],75)
iqr = q3-q1
v_max = q3 + 1.5*iqr
print(v_max)
```


```
0.41000000000000003
```

```
In [21]: print(df[df['f_10']>0.41].shape)
```

```
(29, 50)
```

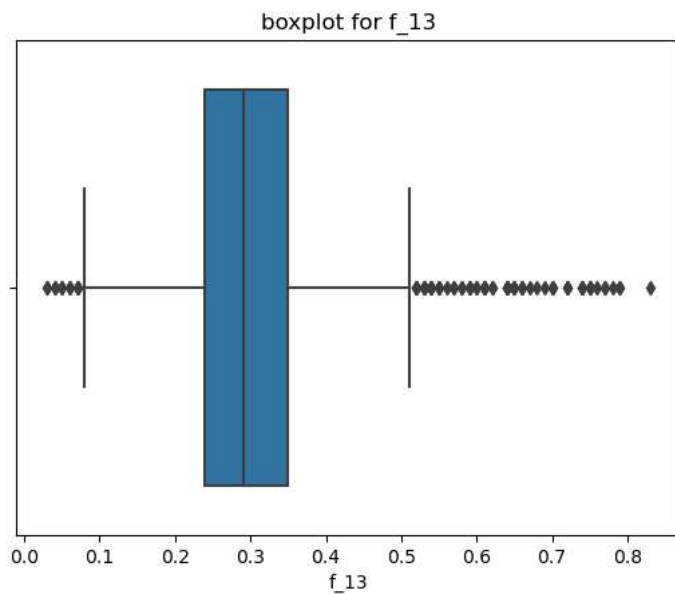
```
In [22]: df['f_10'] = np.where(df['f_10']>0.41,0.41,df['f_10'])
sns.boxplot(x = df['f_10'])
```

Out[22]: <AxesSubplot:xlabel='f\_10'>



```
In [23]: sns.boxplot(df['f_13'])
plt.title('boxplot for f_13')
```

Out[23]: Text(0.5, 1.0, 'boxplot for f\_13')




```
In [24]: q1 = np.percentile(df['f_13'],25)
q3 = np.percentile(df['f_13'],75)
iqr = q3-q1
v_max = q3 + 1.5*iqr
v_min = q3 - 1.5*iqr
print(v_max)
print(v_min)
```

0.5149999999999999  
0.185




```
In [25]: df['f_13'] = np.where(df['f_13']>0.5149999999999999,0.5149999999999999,df['f_13'])
df['f_13'] = np.where(df['f_13']<0.185,0.185,df['f_13'])
sns.boxplot(x = df['f_13'])
```

Out[25]: <AxesSubplot:xlabel='f\_13'>



```
In [26]: sns.boxplot(df['f_44'])
plt.title('boxplot for f_44')
```

Out[26]: Text(0.5, 1.0, 'boxplot for f\_44')



```
In [27]: q1 = np.percentile(df['f_44'],25)
q3 = np.percentile(df['f_44'],75)
iqr = q3-q1
v_max = q3 + 1.5*iqr
v_min = q3 - 1.5*iqr
print(v_max)
```


239.345

```
In [28]: print(df[df['f_44']>239.354].shape)
```

(64, 50)


```
In [29]: df['f_44'] = np.where(df['f_44']>239.345,239.345,df['f_44'])
sns.boxplot(x = df['f_44'])
```

Out[29]: <AxesSubplot:xlabel='f\_44'>



```
In [30]: sns.boxplot(df['f_47'])
```

Out[30]: <AxesSubplot:xlabel='f\_47'>



```
In [31]: q1 = np.percentile(df['f_47'],25)
q3 = np.percentile(df['f_47'],75)
iqr = q3-q1
v_max = q3 + 1.5*iqr
v_min = q3 - 1.5*iqr
print(v_max)
```


18163.97

```
In [32]: print(df[df['f_47']>18163.97].shape)
```

(72, 50)


```
In [33]: df['f_47'] = np.where(df['f_47']>18163.97,18163.97,df['f_47'])  
sns.boxplot(x = df['f_47'])
```

Out[33]: <AxesSubplot:xlabel='f\_47'>



```
In [34]: plt.figure(figsize= (20,20))
corr = df.corr()
corr = corr[abs(corr)>0.9]
sns.heatmap(corr , annot = True)
```

Out[34]: <AxesSubplot:>



Here dropping highly correlated columns are following

f\_9

f\_25

f\_30

f\_48

f\_49

```
In [35]: df.drop(['f_9', 'f_25', 'f_30', 'f_48', 'f_49'], axis = 1, inplace = True)
```

```
In [36]: df.columns
```

```
Out[36]: Index(['f_1', 'f_2', 'f_3', 'f_4', 'f_5', 'f_6', 'f_7', 'f_8', 'f_10', 'f_11',
              'f_12', 'f_13', 'f_14', 'f_15', 'f_16', 'f_17', 'f_18', 'f_19', 'f_20',
              'f_21', 'f_22', 'f_23', 'f_24', 'f_26', 'f_27', 'f_28', 'f_29', 'f_31',
              'f_32', 'f_33', 'f_34', 'f_35', 'f_36', 'f_37', 'f_38', 'f_39', 'f_40',
              'f_41', 'f_42', 'f_43', 'f_44', 'f_45', 'f_46', 'f_47', 'target'],
              dtype='object')
```

```
In [37]: df.shape
```

```
Out[37]: (937, 45)
```

```
In [38]: x = df.drop(['target'], axis = 1)
         y = df['target']
         print(type(x))
         print(type(y))
         print(x.shape)
         print(y.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
(937, 44)
(937,)
```

```
In [39]: from sklearn.model_selection import train_test_split
```

```
In [40]: x_train, x_test, y_train, y_test = train_test_split(x,y , test_size=0.25, random_state= 55)
         print(x_train.shape)
         print(x_test.shape)
         print(y_train.shape)
         print(y_test.shape)
```

```
(702, 44)
(235, 44)
(702,)
(235,)
```

```
In [41]: from sklearn.preprocessing import StandardScaler
```

```
In [42]: scaler = StandardScaler()
         #fit the scaler to train set, it will learn the parameters
         scaler.fit(x_train)
         #transform the train and test sets
         x_train_scaled = scaler.transform(x_train)
         x_test_scaled = scaler.transform(x_test)
```

```
In [43]: x_train_scaled = pd.DataFrame(x_train_scaled, columns = x_train.columns)
         x_test_scaled = pd.DataFrame(x_test_scaled, columns = x_test.columns)
```

```
In [44]: x_train_scaled.head(3)
```

```
Out[44]:
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_10	f_11	...	f_38	f_39	f_40	f_41
0	-0.838841	-0.168880	-1.050269	-0.483907	-0.350673	-0.166102	0.802541	-0.008039	-0.634880	-0.339192	...	-0.262874	-0.429432	-0.779191	0.132371
1	2.006927	-0.123646	1.071601	-0.031226	-0.219135	-0.147670	-0.987585	-1.084454	-0.508013	-0.377937	...	0.013008	-1.256487	-1.583377	-0.041428
2	0.711943	-0.172499	-0.499858	-1.026330	0.197400	-0.181362	0.796218	1.326837	0.380061	-0.091565	...	-0.588407	-0.429432	-0.779191	-0.400611

3 rows × 44 columns

```
In [45]: x_test_scaled.head(3)
```

```
Out[45]:
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_10	f_11	...	f_38	f_39	f_40	f_41
0	-0.902790	-0.147168	-1.007065	-0.532250	1.907387	-0.074541	2.103445	1.095742	-0.508013	0.231867	...	-0.010577	0.489517	0.902287	0.360100
1	1.895015	-0.175515	-1.093557	-0.982231	0.416629	-0.194079	1.195341	0.986276	-0.127410	-0.142101	...	-0.687091	-0.429432	-0.779191	-0.587469
2	-0.087430	-0.117012	1.475213	0.830962	-0.789131	-0.139035	-0.631140	0.475435	1.014399	-0.576712	...	0.495877	-0.613222	-0.413652	-0.004108

3 rows × 44 columns

```
In [46]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [47]: def report(y_test, ypred):
          print('Confusion Matrix\n',confusion_matrix(y_test, ypred))
          print('Classification Report\n',classification_report(y_test, ypred))
          print('Accuracy Score\n',accuracy_score(y_test, ypred))
          def score(model):
              print('Train Score', model.score(x_train, y_train))
              print('Test Score', model.score(x_test, y_test))
```

## LinearRegression

```
In [48]: from sklearn.linear_model import LogisticRegression
```

```
In [49]: model_lr = LogisticRegression(max_iter = 500, solver='liblinear')
          model_lr.fit(x_train,y_train)
```

```
Out[49]: LogisticRegression(max_iter=500, solver='liblinear')
```

```
In [50]: ypred_lr = model_lr.predict(x_test)
          print(ypred_lr[:5])
```

```
[0 0 0 0 0]
```

```
In [51]: report(y_test, ypred_lr)
```

Confusion Matrix

```
[[219  4]
```

```
[ 9  3]]
```

Classification Report

	precision	recall	f1-score	support
0	0.96	0.98	0.97	223
1	0.43	0.25	0.32	12
accuracy			0.94	235
macro avg	0.69	0.62	0.64	235
weighted avg	0.93	0.94	0.94	235

Accuracy Score

```
0.9446808510638298
```

```
In [52]: score(model_lr)
```

Train Score 0.9672364672364673

Test Score 0.9446808510638298

## KNN Classifier

```
In [53]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [54]: model_kn = KNeighborsClassifier(n_neighbors=10)
          model_kn.fit(x_train,y_train)
```

```
Out[54]: KNeighborsClassifier(n_neighbors=10)
```

```
In [55]: ypred_kn = model_kn.predict(x_test)
          print(ypred_kn[:5])
```

```
[0 0 0 0 0]
```

In [56]: `report(y_test, ypred_kn)`

```
Confusion Matrix
[[223  0]
 [ 12  0]]
Classification Report
              precision    recall  f1-score   support

     0       0.95      1.00      0.97       223
     1       0.00      0.00      0.00        12

 accuracy      0.95      0.95      0.95       235
 macro avg     0.47      0.50      0.49       235
 weighted avg   0.90      0.95      0.92       235

Accuracy Score
0.948936170212766
```

In [57]: `score(model_kn)`

```
Train Score 0.9586894586894587
Test Score 0.948936170212766
```

## DecisionTreeClassifier

In [58]: `from sklearn.tree import DecisionTreeClassifier`

In [59]: `model_dt = DecisionTreeClassifier(criterion='entropy', max_depth=4, min_samples_split=25, random_state=45)`  
`model_dt.fit(x_train,y_train)`

Out[59]: `DecisionTreeClassifier(criterion='entropy', max_depth=4, min_samples_split=25, random_state=45)`

In [60]: `ypred_dt = model_dt.predict(x_test)`  
`print(ypred_dt[:5])`

```
[0 0 0 0 0]
```

In [61]: `report(y_test,ypred_dt)`

```
Confusion Matrix
[[215  8]
 [ 8  4]]
Classification Report
              precision    recall  f1-score   support

     0       0.96      0.96      0.96       223
     1       0.33      0.33      0.33        12

 accuracy      0.93      0.93      0.93       235
 macro avg     0.65      0.65      0.65       235
 weighted avg   0.93      0.93      0.93       235

Accuracy Score
0.9319148936170213
```

In [62]: `#Overfit model`  
`score(model_dt)`

```
Train Score 0.98005698005698
Test Score 0.9319148936170213
```

## RandomForestClassifier

In [63]: `from sklearn.ensemble import RandomForestClassifier`

In [64]: `model_rf = RandomForestClassifier(n_estimators=100, criterion='entropy', min_samples_split=25, random_state=45)`  
`model_rf.fit(x_train,y_train)`

Out[64]: `RandomForestClassifier(criterion='entropy', min_samples_split=25, random_state=45)`

In [65]: `ypred_rf = model_rf.predict(x_test)`  
`print(ypred_rf[:5])`

```
[0 0 0 0 0]
```

In [66]: `report(y_test, ypred_rf)`

```
Confusion Matrix
[[222  1]
 [ 12  0]]
Classification Report
              precision    recall  f1-score   support

     0       0.95      1.00      0.97       223
     1       0.00      0.00      0.00        12

 accuracy          0.94       235
 macro avg         0.47       235
 weighted avg      0.90       235

Accuracy Score
0.9446808510638298
```

In [67]: `score(model_rf)`

```
Train Score 0.9629629629629629
Test Score 0.9446808510638298
```

## BaggingClassifier

In [68]: `from sklearn.ensemble import BaggingClassifier`

In [69]: `model_bg = BaggingClassifier(base_estimator=model_dt, n_estimators=100, max_samples = x_train.shape[0],  
max_features=x_train.shape[1], random_state=45)  
model_bg.fit(x_train, y_train)`

Out[69]: `BaggingClassifier(base_estimator=DecisionTreeClassifier(criterion='entropy',  
max_depth=4,  
min_samples_split=25,  
random_state=45),  
max_features=44, max_samples=702, n_estimators=100,  
random_state=45)`

In [70]: `ypred_bg = model_bg.predict(x_test)  
print(ypred_bg[:5])`

```
[0 0 0 0 0]
```

In [71]: `report(y_test, ypred_bg)`

```
Confusion Matrix
[[219  4]
 [  6  6]]
Classification Report
              precision    recall  f1-score   support

     0       0.97      0.98      0.98       223
     1       0.60      0.50      0.55        12

 accuracy          0.96       235
 macro avg         0.79       235
 weighted avg      0.95       235

Accuracy Score
0.9574468085106383
```

In [72]: `score(model_bg)`

```
Train Score 0.9729344729344729
Test Score 0.9574468085106383
```

## AdaBoostClassifier

In [73]: `from sklearn.ensemble import AdaBoostClassifier`



```
In [74]: model_ab = AdaBoostClassifier(base_estimator = model_dt, n_estimators = 90, random_state=45)
model_ab.fit(x_train,y_train)
```

```
Out[74]: AdaBoostClassifier(base_estimator=DecisionTreeClassifier(criterion='entropy',
max_depth=4,
min_samples_split=25,
random_state=45),
n_estimators=90, random_state=45)
```

```
In [75]: ypred_ab = model_ab.predict(x_test)
print(ypred_ab[:5])
```

```
[0 0 0 0 0]
```

```
In [76]: report(y_test, ypred_ab)
```

```
Confusion Matrix
[[220  3]
 [ 7  5]]
Classification Report
precision    recall  f1-score   support

     0       0.97     0.99     0.98       223
     1       0.62     0.42     0.50        12

 accuracy          0.96       235
 macro avg       0.80       0.70       0.74       235
 weighted avg    0.95       0.96       0.95       235

Accuracy Score
0.9574468085106383
```

```
In [77]: score(model_ab)
```

```
Train Score 1.0
Test Score 0.9574468085106383
```

**From the above we can conclude that best model is of BaggingClassifier**

```
In [78]: import pickle
```

```
In [79]: pickle.dump(model_bg,open('model_bg.pkl','wb'))
```

```
In [80]: model = pickle.load(open('model_bg.pkl','rb'))
```

```
In [81]: data = x_train.sample(20)
print(data)
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_10	f_11	\
511	17	236	256.73	1301.15	119	331875	25.98	12.110	0.41	81.5	
413	152	102	839.88	1231.77	90	255000	38.96	7.800	0.20	101.6	
346	59	55	1407.87	913.91	94	137500	41.25	6.510	0.16	64.6	
642	74	11	65.00	386.27	105	89100	57.45	9.610	0.17	107.2	
425	177	59	844.44	1106.71	101	147500	40.80	7.730	0.19	77.2	
18	8	1035	1304.90	1281.72	68	2587500	35.90	7.510	0.21	167.8	
36	26	81	1705.68	1492.58	19	202500	35.65	6.320	0.18	81.7	
225	68	94	1062.21	719.81	83	235000	31.88	7.670	0.24	79.4	
331	41	204	906.97	960.46	107	510000	39.53	7.190	0.18	88.5	
764	32	145	20.19	418.42	72	1174500	61.51	16.775	0.28	184.6	
184	179	74	1751.49	1095.38	119	185000	39.46	8.790	0.22	82.6	
259	192	162	1178.77	993.60	99	405000	34.73	6.880	0.20	93.3	
634	66	10	76.50	441.70	143	81000	49.50	16.775	0.34	104.1	
97	87	125	1785.22	1487.77	3	312500	38.87	10.310	0.27	114.1	
92	82	83	1260.70	1474.01	66	207500	39.83	7.830	0.20	81.7	
729	82	12	13.33	483.92	34	97200	73.33	5.470	0.07	112.0	
539	45	65	511.22	1709.25	144	91406	24.95	9.470	0.38	43.2	
5	6	54	1438.13	544.91	82	135000	44.67	6.920	0.15	86.0	
459	27	91	982.67	899.04	49	127968	27.27	9.120	0.33	54.0	
873	141	10	52.70	353.40	95	81000	55.10	7.200	0.13	95.6	

	...	f_38	f_39	f_40	f_41	f_42	f_43	f_44	f_45	f_46	\
511	...	49.91	143	86	1544.34	270.42	156.69	67.520	9.86	1	
413	...	24.71	85	63	800.00	500.00	262.50	117.620	3.05	0	
346	...	33.00	85	63	651.92	500.00	138.46	113.930	4.71	0	
642	...	7.75	99	67	284.60	180.00	150.00	51.960	1.90	0	
425	...	24.73	85	63	509.90	304.14	208.37	45.110	2.45	0	
18	...	91.89	78	55	4419.56	1234.91	632.82	239.345	6.98	0	
36	...	30.37	78	55	901.39	269.26	178.33	59.630	5.05	0	
225	...	37.28	64	39	912.41	412.31	196.18	128.280	4.65	0	
331	...	65.05	85	63	1408.01	522.02	288.58	116.270	4.88	0	
764	...	34.46	82	50	1793.24	742.16	445.58	100.170	4.02	0	
184	...	27.12	78	55	694.62	291.55	192.72	51.620	3.60	0	
259	...	46.51	64	39	1265.90	509.90	227.86	140.980	5.56	0	
634	...	7.47	99	67	254.56	0.00	0.00	0.000	0.00	1	
97	...	24.02	78	55	901.39	500.00	288.89	135.640	3.12	0	
92	...	31.09	78	55	1051.19	316.23	135.59	80.160	7.75	0	
729	...	7.75	102	73	402.49	127.28	95.46	63.640	4.22	0	
539	...	49.01	143	86	0.00	0.00	0.00	0.000	0.00	0	
5	...	18.26	89	69	608.28	200.00	150.00	52.220	4.06	0	
459	...	43.89	133	85	850.18	265.17	96.93	90.130	8.77	0	
873	...	8.86	82	50	360.00	90.00	67.50	45.000	5.33	0	

	f_47
511	8424.77
413	5147.43
346	4612.24
642	8942.39
425	4065.36
18	6165.56
36	3378.60
225	5066.66
331	3009.93
764	2326.02
184	5158.25
259	5606.99
634	7133.13
97	4736.75
92	7584.28
729	5509.43
539	8487.58
5	18163.97
459	12460.92
873	2859.39

[20 rows x 44 columns]

```
In [82]: model_pred = model.predict(data)
print(model_pred)
print(len(model_pred))
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0]
20
```